

Instruction and Return Barriers

Code Integrity and Execution Correctness

Why?

- JIT performance
 - Profile based
 - Tracing
 - Speculation
 - Garbage collection
 - Managed Runtime Resource Reclamation
- JIT compilation costly
 - Compile as little as possible

Scenarios

- Instructions ahead are invalid
- Invalid Frame Data
 - Frame part of execution
 - Barrier may trigger GC/heap barrier
- Execution Must Be Stopped
- ?

Instructions ahead are invalid

- Stale Instructions
 - I/D Cache Inconsistency
 - Invalid Materialized Constants
 - Environmental Invalidation
 - Ineffective/Unwanted Instructions
- Missing instructions
 - Incomplete Compilation
 - Tracing JIT
 - Unavailable Data
 - Lazy loading
 - Linking without all referred object files

Invalid Frame Data

- Invalid
 - Stale references
 - Reference refer to moved objects
- Missing
 - Not mounted on native stack
 - Continuations

Execution Stoppage

- Exclusive access
 - E.g. object graph
- Quiescent state
 - As in RCU/QSBR
 - E.g. a JIT:ed method is reclaimed after a QS

Implementation Strategies

- Explicit Branching
 - Conditional jumps to stub routines.
- Trapping Instructions
 - Deliberately causing hardware traps (e.g., illegal instructions) to invoke slowpaths.
- Self- or Cross-Modifying Code
 - Patching instructions at runtime to redirect control flow.
- Destination Changes
 - Altering jump targets.
- Return Address Patching
 - Modifying the stack return address to hijack control flow upon function return.

Placement

- Function Entry
 - Entry Barrier, Entry Guard, Receiver Check
- Function Exit
 - Callee, Caller, Return address
- Guts
 - Deoptimization traps
 - E.g. Dead Code Resurrection
 - Execution Guards
 - Striped mined loops
 - Initialization Guards
 - Once barriers